

loan_prediction_problem

September 2, 2020

In [this dataset](#), we need to predict whether or not to approve a loan based on the past information of the person. This is a classification problem and we will use machine learning, Decision Tree Classifier model, to make the prediction.

1 Import Libraries

First, we import necessary libraries, such as:

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set()

import warnings
warnings.filterwarnings('ignore')
```

2 Import The Data

```
[2]: train = pd.read_csv('/kaggle/input/loan-prediction-problem-dataset/
↳train_u6lujuX_CVtuZ9i.csv')
```

3 Read The Data

First, let's see the first 5 rows to familiarize ourselves with the data.

```
[3]: train.head()
```

```
[3]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	

3	LP001006	Male	Yes	0	Not Graduate	No
4	LP001008	Male	No	0	Graduate	No

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

To get more details, we are going to print `info()` and `describe()` to make a quick observation and gain some insight from it.

```
[4]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                614 non-null    object
1   Gender                 601 non-null    object
2   Married                611 non-null    object
3   Dependents             599 non-null    object
4   Education              614 non-null    object
5   Self_Employed          582 non-null    object
6   ApplicantIncome        614 non-null    int64
7   CoapplicantIncome      614 non-null    float64
8   LoanAmount             592 non-null    float64
9   Loan_Amount_Term       600 non-null    float64
10  Credit_History         564 non-null    float64
11  Property_Area          614 non-null    object
12  Loan_Status            614 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
```

```
[5]: train.describe(include='all')
```

```
[5]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
count	614	601	611	599	614	582	

unique	614	2	2	4	2	2
top	LP002239	Male	Yes	0	Graduate	No
freq	1	489	398	345	480	500
mean	NaN	NaN	NaN	NaN	NaN	NaN
std	NaN	NaN	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN	NaN	NaN

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term \
count	614.000000	614.000000	592.000000	600.00000
unique	NaN	NaN	NaN	NaN
top	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN
mean	5403.459283	1621.245798	146.412162	342.00000
std	6109.041673	2926.248369	85.587325	65.12041
min	150.000000	0.000000	9.000000	12.00000
25%	2877.500000	0.000000	100.000000	360.00000
50%	3812.500000	1188.500000	128.000000	360.00000
75%	5795.000000	2297.250000	168.000000	360.00000
max	81000.000000	41667.000000	700.000000	480.00000

	Credit_History	Property_Area	Loan_Status
count	564.000000	614	614
unique	NaN	3	2
top	NaN	Semiurban	Y
freq	NaN	233	422
mean	0.842199	NaN	NaN
std	0.364878	NaN	NaN
min	0.000000	NaN	NaN
25%	1.000000	NaN	NaN
50%	1.000000	NaN	NaN
75%	1.000000	NaN	NaN
max	1.000000	NaN	NaN

```
[6]: train.isnull().sum().sort_values(ascending=False)
```

```
[6]: Credit_History      50
      Self_Employed      32
      LoanAmount          22
      Dependents          15
      Loan_Amount_Term    14
      Gender              13
      Married              3
      Loan_Status          0
```

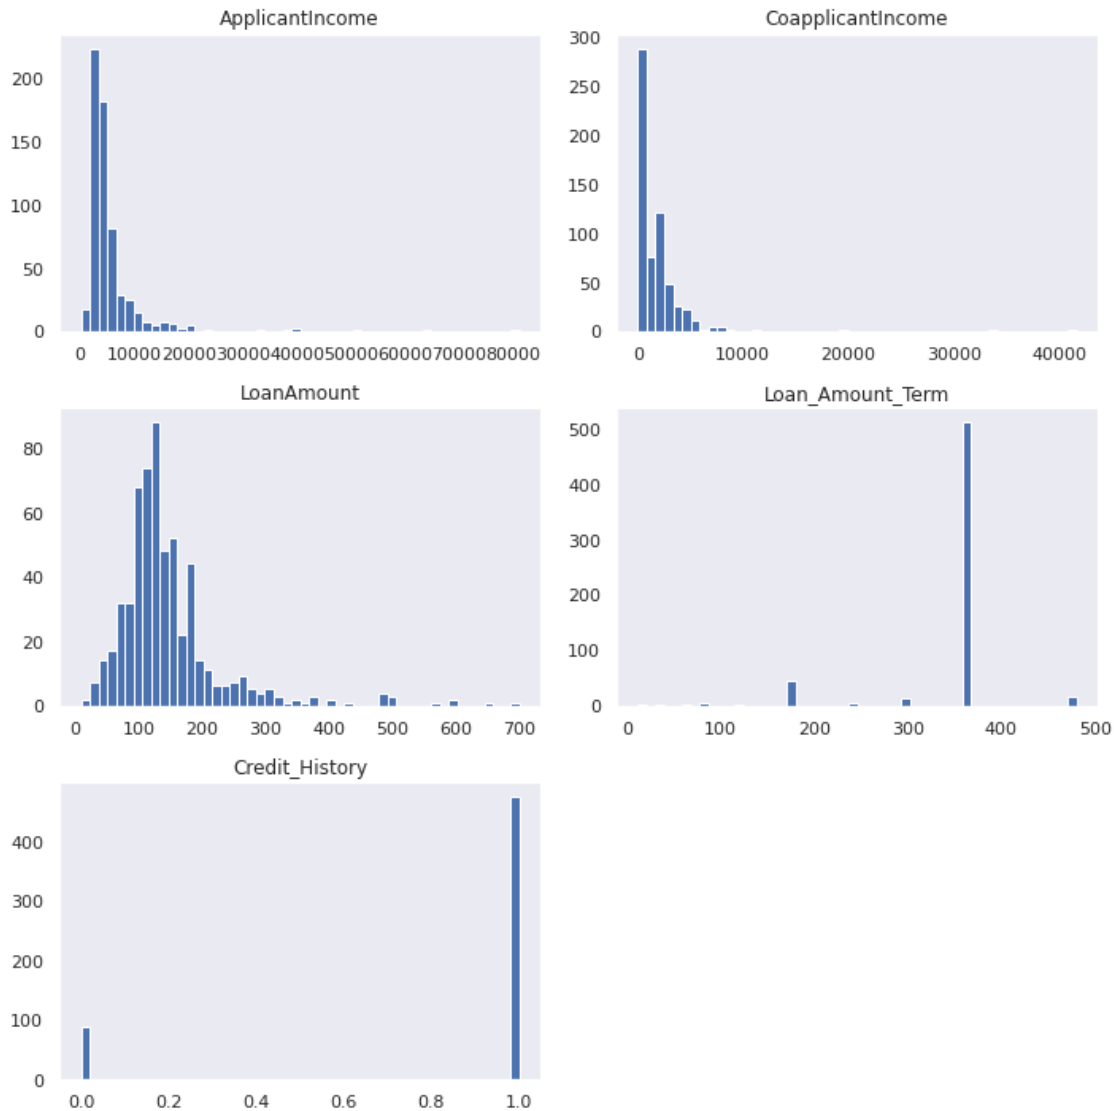
```
Property_Area      0
CoapplicantIncome  0
ApplicantIncome    0
Education          0
Loan_ID            0
dtype: int64
```

3.0.1 Quick observation on the combined data

- Total loaner: 614
- Feature that can be dropped from training immediately:
 - **Loan_ID**
- The **Loan_Status** feature, as target array, can be used as a reference to fill missing values, so we will not drop it immediately.
- The **Dependents** feature is given in categorical but contain numerical variables. Therefore, we have to converted it to numerical variables.
- The **Credit_History** feature is given in numerical, with 1 means 'Yes' and 0 means 'No'. We will converted it to categorical feature.
- Features that have missing values:
 - **Credit_History**: 50
 - **Self_Employed**: 32
 - **LoanAmount**: 22
 - **Dependents**: 15
 - **Gender**: 14
 - **Loan_Amount_Term**: 13
 - **Married**: 3

3.0.2 Plot The Distribution of Numerical Features

```
[7]: #plot the distribution of numerical features
train.hist(bins=50,figsize=(10,10),grid=False)
plt.tight_layout()
plt.show()
```



We can see we got right-skewed and left-skewed. We will fix this in the next step by taking the log of the values to make it normally distributed. By making it normally distributed, we can improve our model.

4 Exploratory Data Analysis

4.0.1 Drop Features

```
[8]: #drop feature
train.drop(['Loan_ID'], axis=1, inplace=True)
```

4.0.2 Change to Numerical

```
[9]: #check unique values
train['Dependents'].unique()
```

```
[9]: array(['0', '1', '2', '3+', nan], dtype=object)
```

```
[10]: #replace '3+' with '3'
train['Dependents'].replace('3+', '3', inplace=True)

#change to numerical
train['Dependents'] = train['Dependents'].astype('float')
```

4.0.3 Change to Categorical

```
[11]: #check unique values
train['Credit_History'].unique()
```

```
[11]: array([ 1.,  0., nan])
```

```
[12]: #replace 1.0 with 'Y' and 0.0 with 'N'
train['Credit_History'].replace({1.: 'Y', 0.: 'N'}, inplace=True)

#change to categorical
train['Credit_History'] = train['Credit_History'].astype('object')
```

4.0.4 Fill Missing Value:

- Fill with mode()

```
[13]: #fill missing values with mode
features_fill_with_mode = ['Self_Employed',
                           'Dependents',
                           'Gender',
                           'Loan_Amount_Term']

for feature in features_fill_with_mode:
    train[feature].fillna(train[feature].mode()[0], inplace=True)
```

- Fill with mean()

```
[14]: #fill missing values with mean
train['LoanAmount'].fillna(train['LoanAmount'].mean(), inplace=True)
```

- Credit_History Feature

Before we fill missing values in `Credit_History` feature, we will take a deeper look by plotting it.

```
[15]: sns.countplot(x='Credit_History', hue='Loan_Status', data=train);
```



From the plot above, we can see that `Credit_History` is important feature. Most people with 0 credit history didn't get a loan. But, most people who got credit history have so much better chance to get a loan.

Since `Credit_History = 'Y'` is the value that appears most often in both `Loan_Status`, so we will fill missing values with 'Y'

```
[16]: train['Credit_History'].fillna('Y', inplace=True)
```

- Married Feature

For start, we will check if the missing values in the `Married` feature have `Dependents` or `CoapplicantIncome` more than 0, and fill it with 'Yes' if true and 'No' if otherwise.

```
[17]: #check Dependents and CoapplicantIncome
mask = ((train['Dependents'] > 0) | (train['CoapplicantIncome'] > 0)) \
        & \
        train['Married'].isnull()

train[mask][['Married', 'Dependents', 'CoapplicantIncome']]
```

```
[17]:    Married  Dependents  CoapplicantIncome
      104      NaN        0.0              754.0
```

```
[18]: #Fill missing values
      train.loc[mask, 'Married'] = 'Yes'
      train['Married'].fillna('No', inplace=True)
```

4.0.5 Target Array

let's look at the target distribution


```
[19]: sns.countplot(train['Loan_Status']);
```



From the distribution above, we can consider that the data is not imbalanced. So, we can straight to the next step: change it to numerical feature.

```
[20]: #transform to numerical
      train['Loan_Status'] = train['Loan_Status'].apply(lambda x: 1 if x=='Y' else 0)

      #correlation
      sns.heatmap(train.corr(), annot=True);
```

```
[21]: #copy
target_array = train['Loan_Status'].copy()

#drop
train.drop(['Loan_Status'], axis=1, inplace=True)
```

4.0.6 Creating new features

```
[22]: #create total income feature
train['Total_Income'] = train['ApplicantIncome'] + train['CoapplicantIncome']

#create average loan amount feature (per day)
train['Loan_Amount_Avg'] = train['LoanAmount'] / train['Loan_Amount_Term']

#drop
train.drop(['ApplicantIncome', 'CoapplicantIncome'], axis=1, inplace=True)
```

4.0.7 Epilogue

- Check for any missing values

```
[23]: #missing values
print(train.isnull().any().sum())
```

0

- Normality Test

```
[24]: #define a normality test function
def normalityTest(data, alpha=0.05):
    """data (array) : The array containing the sample to be tested.
        alpha (float) : Significance level.
        return True if data is normal distributed"""

    from scipy import stats

    statistic, p_value = stats.normaltest(data)

    #null hypothesis: array comes from a normal distribution
    if p_value < alpha:
        #The null hypothesis can be rejected
        is_normal_dist = False
    else:
        #The null hypothesis cannot be rejected
        is_normal_dist = True

    return is_normal_dist
```

```
[25]: #check normality of all numericaal features and transform it if not normal,
      ↪distributed
for feature in train.columns:
    if (train[feature].dtype != 'object'):
        if normalityTest(train[feature]) == False:
            train[feature] = np.log1p(train[feature])
```

- Creating Dummies

```
[26]: #create dummies
train = pd.get_dummies(train, drop_first=True)

print(train.shape)
display(train.head())
```

(614, 12)

Dependents LoanAmount Loan_Amount_Term Total_Income Loan_Amount_Avg \

0	0.000000	4.993232	5.888878	8.674197	0.341247
1	0.693147	4.859812	5.888878	8.714732	0.304211
2	0.000000	4.204693	5.888878	8.006701	0.168335
3	0.000000	4.795791	5.888878	8.505525	0.287682
4	0.000000	4.955827	5.888878	8.699681	0.330502

	Gender_Male	Married_Yes	Education_Not Graduate	Self_Employed_Yes	\
0	1	0	0	0	
1	1	1	0	0	
2	1	1	0	1	
3	1	1	1	0	
4	1	0	0	0	

	Credit_History_Y	Property_Area_Semiurban	Property_Area_Urban
0	1	0	1
1	1	0	0
2	1	0	1
3	1	0	1
4	1	0	1

- Creating features matrix (X) and target array (y)

```
[27]: X = train
      y = target_array
```

5 Creating a Model

We begin by splitting data into two subsets: for training data and for testing data.

```
[28]: from sklearn.model_selection import train_test_split

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
      ↪random_state = 0)
```

Model training : Decision Tree Classifier

```
[29]: from sklearn.tree import DecisionTreeClassifier

      #create a model
      model = DecisionTreeClassifier()
```

```
[30]: #search grid for optimal parameters
      from sklearn.model_selection import GridSearchCV

      param_grid = {'random_state' : [0,42],
                    'max_depth': [1,10,100]}
```

```

grid = GridSearchCV(model, param_grid, cv=5)

grid.fit(X_train, y_train)

print(grid.best_params_)
print(grid.best_score_)

```

```

{'max_depth': 1, 'random_state': 0}
0.8028504260946224

```

```
[31]: from sklearn.metrics import classification_report
```

```

#use the best model
model = grid.best_estimator_

#make a prediction
y_predict = model.predict(X_test)

#calculate classification report
print(classification_report(y_test,y_predict))

```

	precision	recall	f1-score	support
0	0.92	0.41	0.57	58
1	0.81	0.99	0.89	145
accuracy			0.82	203
macro avg	0.87	0.70	0.73	203
weighted avg	0.84	0.82	0.80	203

```
[ ]:
```